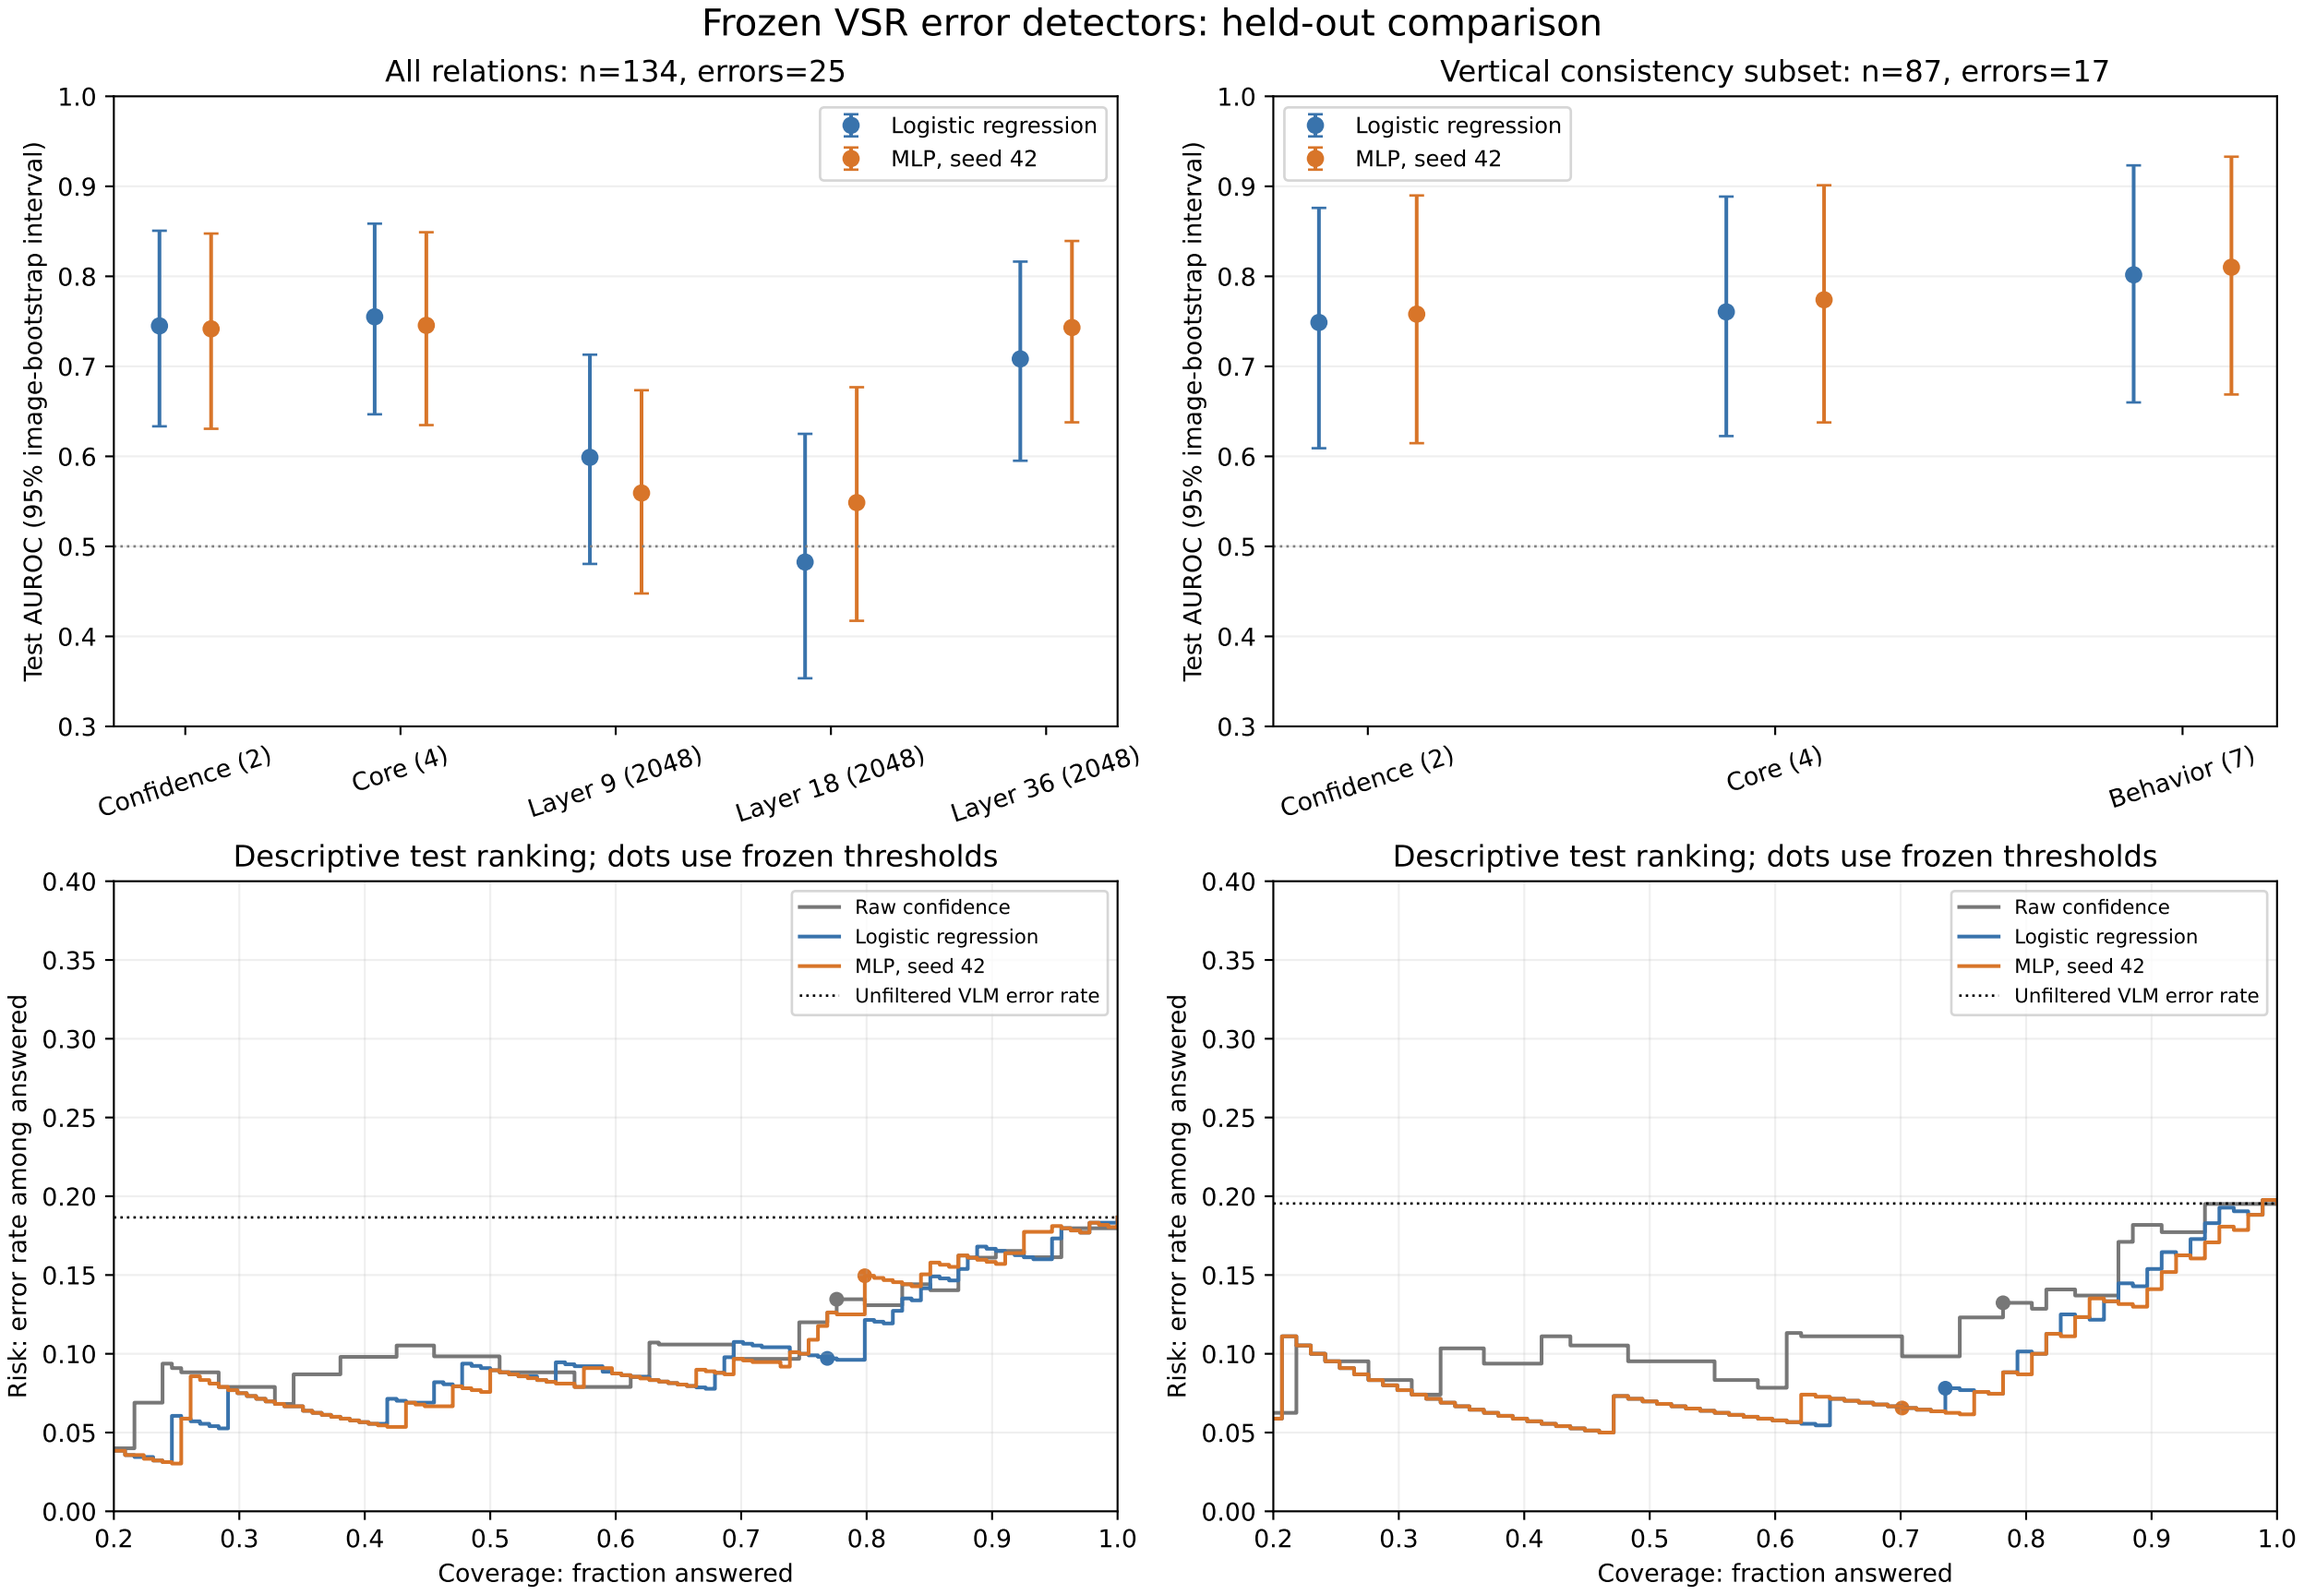
Frozen VSR error detectors: held-out comparison
All relations: n=134, errors=25
Vertical consistency subset: n=87, errors=17
Test AUROC (95% image-bootstrap interval)
Logistic regression
MLP, seed 42
Confidence (2)
Core (4)
Layer 9 (2048)
Layer 18 (2048)
Layer 36 (2048)
Behavior (7)
Descriptive test ranking; dots use frozen thresholds
Raw confidence
Unfiltered VLM error rate
Risk: error rate among answered
Coverage: fraction answered